

Daniel Zou

☎ (703) 939-0901 | ✉ dlzou@berkeley.edu | 🌐 dlzou | 🐙 dlzou | 📄 dlzou.github.io

EDUCATION

University of California, Berkeley

Berkeley, CA

B.Sc. Electrical Engineering and Computer Sciences, with Honors; GPA: 3.94/4.00

Aug 2019 – May 2023

- Selected coursework: OS, databases, algorithms, compilers, security, ML, DL

Academy of Science and Technology

The Woodlands, TX

High School Diploma; Rank: 2/712

Aug 2015 – May 2019

RESEARCH EXPERIENCE

Berkeley Sky Computing Lab

Berkeley, CA

Undergraduate Research Assistant

Sep 2022 – Present

- Contributing to Alpa, a system for automatically parallelizing the training and serving of large-scale neural networks.
- Implementing scheduled swapping between GPU and main memories in the Alpa runtime to enable deployment of large models on clusters with limited GPU resources.
- Supervised by Lianmin Zheng and Professor Joseph E. Gonzalez.

Berkeley RISE Lab

Berkeley, CA

Undergraduate Research Assistant

Mar 2021 – Present

- Contributed to NumS, a high performance distributed numerical library for Python with a NumPy-like interface, with support for Ray and Dask backends.
- Studied and innovated on techniques like operator fusion and cost estimation to dynamically schedule sparse-aware tensor operations on multi-node systems. Publication is in preparation.
- Implemented distributed algorithms such as quickselect for block-partitioned arrays, then analyzed the scaling properties of said algorithms through benchmarks on an EC2 cluster.
- Supervised by Melih Elibol, PhD and Professor Ion Stoica.

WORK EXPERIENCE

NVIDIA Corporation

Santa Clara, CA

Autonomous Vehicles Software Infrastructure Intern

May 2022 – Aug 2022

- Designed and built on-call automation and metrics collection tooling from the ground up, for use by MagLev, NVIDIA's autonomous vehicle infrastructure organization.
- Architected a concurrent, event-driven service with WebSockets, producer-consumer queue, in-memory cache, SQL database, and a querying command interface.
- Worked with multiple stakeholder representatives to revise features and improve usability of the tooling during development and beta testing. Presented the product to 50+ engineers and managers, then engaged with discussions of expanding product scope to other organizations.
- Deployed for use by 300+ engineers and business partners, and growing.

ORGANIZATIONS

Computer Science Undergraduate Association

Berkeley, CA

VP of Technology & Root Staff

Sep 2020 – Aug 2022

- Administered a GPU cluster, a web hosting server, and other free computing services used by 400+ members.
- Maintained a full stack web application written in Django.
- Configured a multi-instance Postfix mail server.
- Served as the leader of CSUA Root Staff for the Fall 2021 semester upon election to VP of Technology. Acted as liaison with campus IT staff to resolve issues regarding networking services.

Computer Science Mentors

Berkeley, CA

Junior Mentor

Jan 2021 – May 2021

- Taught CS 61C: Great Ideas in Computer Architecture, which covers C, RISC-V assembly, CPU logic design, memory hierarchy, and parallel programming.
- Mentored a group of student every week by giving short lectures and practice problem walk-throughs.

FIRST Robotics Team 1477 (Texas Torque)

The Woodlands, TX

Lead Programmer

Jun 2018 – May 2019

- Recruited a team of 6 programmers and oversaw the development of the robot control system.
- Designed new autonomous control framework that enables scheduling actions for multiple subsystems in parallel. Built a custom object tracking CV pipeline that runs on a coprocessor.
- Reached the quarterfinal of the 2019 FRC World Championship as a team.

SELECTED PROJECTS

PintOS

- An instructional OS written in C and x86. Completed in CS 162: Operation Systems and Systems Programming over the course of a semester.
- Implemented basic user processes, kernel threads, concurrency primitives, and a concurrent filesystem with a buffer cache.

NP-Hard Graph Problem Solver | [GitHub](#)

- An open-ended project to find optimal solutions for a set of inputs to a given NP-hard graph problem. Completed in CS 170: Efficient Algorithms and Intractable Problems.
- Designed a greedy algorithm as a baseline and a simulated annealing algorithm to explore the space of potentially better solutions. Wrote scripts to merge best solutions over multiple runs.

NumC: Toy Reimplementation of Core NumPy Functionality

- Completed in CS 61C: Great Ideas in Computer Architecture.
- Lightweight numerical library for Python, with underlying data structures written in C for performance.
- Supported efficient matrix slicing and common operations like add, abs, matmul, matpow. Optimized matmul with threading, loop tiling, and SIMD intrinsics. Optimized matpow with dynamic programming.

Deep NLP Chatbot | [Blog](#)

- Self-studied RNNs and its variants, as well as the TensorFlow ML framework.
- Implemented an attention-based seq2seq model that uses a beam search decoder for evaluation.

AWARDS

National Merit Finalist (2019): selected among high school seniors scoring in the top 0.5% of the PSAT.

National AP Scholar (2019): received the maximum score of 5 on fourteen Advanced Placement exams.

SKILLS

Programming: Python, C, Go, Java, OCaml, Lua, SQL, RISC-V, x86

Libraries: SciPy, PyTorch, JAX, Ray, Spark, Django

Technologies: Linux, Git, Jira, Docker, AWS EC2

Other: native English and Mandarin Chinese, technical and blog writing